

IoT Data Processing System - Comprehensive Documentation

This document provides a detailed overview of the IoT data processing system, including architecture, component interactions, scaling strategies, and monitoring approaches.

System Architecture

High-Level Architecture

[Architecture Diagram: IoT Devices → IoT Ingest API → Kafka → IoT Consumer Worker → Database]

The system follows a microservices architecture with clear separation of concerns:

- IoT Ingest API:** Receives device registration and telemetry data via REST API
- Apache Kafka:** Acts as a message broker for asynchronous communication
- IoT Consumer Worker:** Processes telemetry messages from Kafka and stores them in the database
- MySQL Database:** Stores device information and telemetry data

Data Flow

- Devices register through the IoT Ingest API
- Registered devices send telemetry data to the ingest API
- The API validates devices and publishes telemetry to Kafka
- Consumer workers read from Kafka, validate devices again, and store data in the database
- Data is available for querying and analysis

Component Details

IoT Ingest API Service

Purpose: Provides REST endpoints for device management and telemetry ingestion.

Key Components:

- DeviceController:** Handles device registration and retrieval
- TelemetryController:** Accepts telemetry data and publishes to Kafka
- DeviceBusinessService:** Business logic for device operations
- TelemetryProducerService:** Kafka producer service

Configuration Details:

```
server.port: 8080
spring.kafka.bootstrap-servers: localhost:9092
spring.datasource.url: jdbc:mysql://localhost:3306/iotdatabase
```

IoT Consumer Worker Service

Purpose: Consumes telemetry messages from Kafka and stores them in the database.

Key Components:

- TelemetryConsumer:** Kafka consumer with batch processing capabilities
- DeviceRepository:** Data access for device validation
- TelemetryRepository:** Data access for storing telemetry records

Configuration Details:

```
server.port: 8081
spring.kafka.bootstrap-servers: localhost:9092
spring.kafka.consumer.group-id: iot-telemetry-group
spring.kafka.consumer.max-poll-records: 500
```

Service Interconnection

How Services Interact

1. Device Registration Flow

- Client sends POST request to `/api/v1/devices` on IoT Ingest API
- API validates and saves device to MySQL database
- Device information is now available for both services

2. Telemetry Ingestion Flow

- Device sends telemetry data to `/api/v1/telemetry` on IoT Ingest API
- API validates that device is registered
- API publishes telemetry data to Kafka topic `iot-telemetry`
- IoT Consumer Worker reads messages from Kafka
- Consumer validates device exists in database
- Consumer saves telemetry data to MySQL

3. Data Consistency

Both services share the same database for device information, ensuring consistency in device validation.

Kafka provides at-least-once delivery semantics, ensuring no telemetry data is lost.

Scaling Strategies

Horizontal Scaling Approaches

IoT Ingest API Scaling

- Deploy multiple instances behind a load balancer
- Stateless design allows easy horizontal scaling
- Consider using API gateway for rate limiting and authentication

IoT Consumer Worker Scaling

- Deploy multiple consumer instances with the same group ID
- Kafka automatically partitions the load across consumers
- Increase Kafka topic partitions to allow more parallel consumers
- Adjust `concurrency` setting for more threads per instance

Database Scaling

- Implement read replicas for reporting queries
- Consider sharding by device ID for very large deployments
- Use connection pooling and optimize queries

Kafka Scaling

- Increase topic partitions as needed
- Monitor consumer lag to identify bottlenecks
- Consider Kafka cluster expansion for high throughput scenarios

Performance Optimization

Component	Optimization Strategy
IoT Ingest API	Async processing, connection pooling, efficient serialization
IoT Consumer Worker	Batch processing, database bulk inserts, connection pooling
Database	Indexing, query optimization, appropriate hardware
Kafka	Proper partitioning, compression, monitoring

Monitoring and Observability

Monitoring Strategy

Application Metrics

- IoT Ingest API** exposes Prometheus metrics via Spring Boot Actuator
- Key metrics: HTTP request rates, error rates, JVM metrics, custom telemetry counters
- IoT Consumer Worker** tracks message processing rates, error rates, and database performance

Infrastructure Monitoring

- Monitor server resources (CPU, memory, disk I/O)
- Database performance metrics (query times, connection pool usage)
- Kafka metrics (topic throughput, consumer lag, broker performance)

Alerting

Set up alerts for:

- High error rates in either service
- High consumer lag in Kafka
- Database connection issues
- Service downtime

Logging

- Structured logging with correlation IDs for tracing requests across services
- Centralized log aggregation (ELK stack, Splunk, etc.)
- Log retention policies for troubleshooting and compliance

Key Performance Indicators

Metric	Description	Target
API Response Time	P95 response time for ingest API	< 200ms
Telemetry Processing Lag	Time between ingestion and storage	< 1 second
Kafka Consumer Lag	Messages waiting to be processed	< 1000 messages
Error Rate	Percentage of failed requests/operations	< 0.1%

Deployment Considerations

Containerization

Both services are containerized with Docker for consistent deployment across environments.

Orchestration

For production deployment, consider using Kubernetes for:

- Automated scaling based on load
- Self-healing capabilities
- Rolling updates with zero downtime
- Resource management

Environment Configuration

Use environment-specific configuration files or config management tools for:

- Database connection strings
- Kafka broker addresses
- External service URLs
- Feature flags

Future Enhancements

- Add real-time analytics capabilities
- Implement device management dashboard
- Add support for more telemetry formats
- Implement data retention policies
- Add anomaly detection for telemetry data